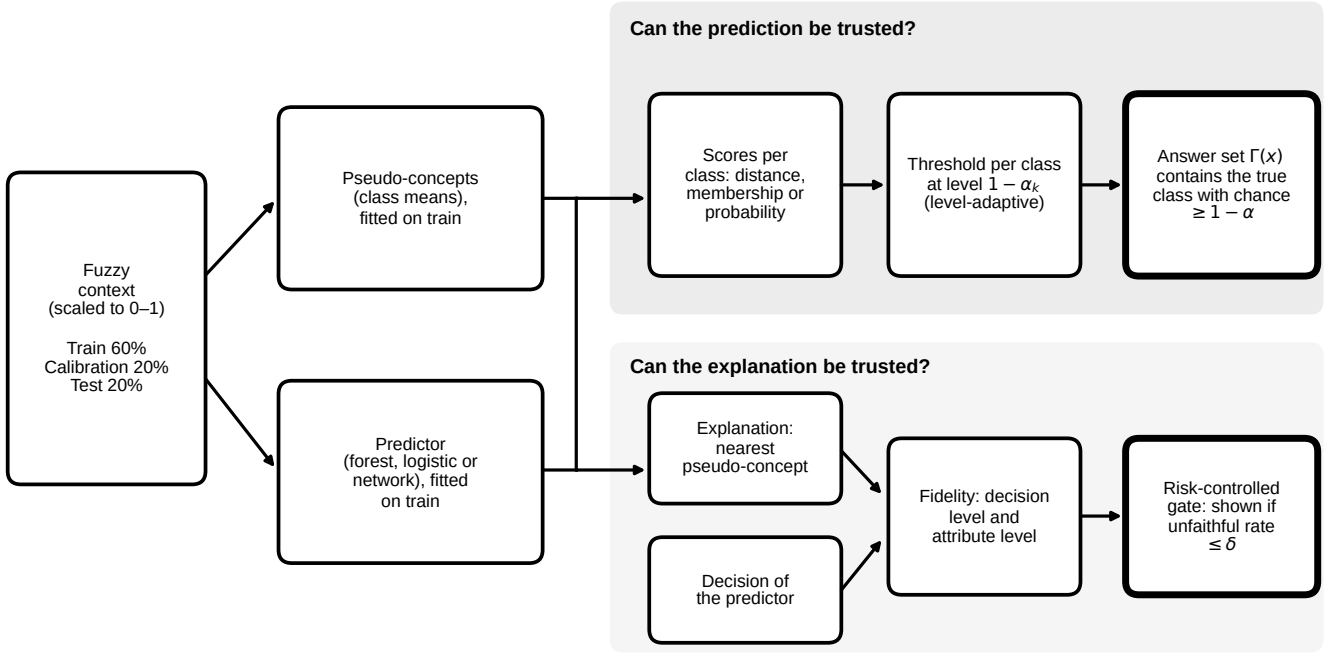
Fuzzy
context
(scaled to 0–1)
Train 60%
Calibration 20%
Test 20%
Pseudo-concepts
(class means),
fitted on train
Predictor
(forest, logistic or
network), fitted
on train
Can the prediction be trusted?
Scores per
class: distance,
membership or
probability
Threshold per class
at level $1 - \alpha_k$
(level-adaptive)
Answer set $\Gamma(x)$
contains the true
class with chance
$\geq 1 - \alpha$
Can the explanation be trusted?
Explanation:
nearest
pseudo-concept
Decision of
the predictor
Fidelity: decision
level and
attribute level
Risk-controlled
gate: shown if
unfaithful rate
$\leq \delta$